

# Introduction to Deep Learning 2022

---

**Lecturers:** Wojtek Kowalczyk  
Andrius Bernatavicius

**TA's:**  
Andrius Bernatavicius (coordinator)  
&  
5 TA's (details will follow soon)

# Course Objectives

---

## Classical Neural Networks:

- *Basics of Machine Learning and Pattern Recognition (?)*
- Perceptron: Linear Separability, Convergence Theorem, Cover's Theorem
- Multi-layer Perceptron and its expressive power
- Algorithms for training MLPs: Stochastic Gradient Descent and its variants; Backpropagation
- Alternative activation and loss functions; Initialization, Regularization, Dropout, Batch Normalization
- Introduction to Keras, TensorFlow, GPU-computing

# Course Objectives



## Deep Neural Networks:

- Convolutional Networks; key architectures and applications; Transfer Learning
- Recurrent Neural Networks: Backpropagation through time, LSTM and GRU Networks
- Attention mechanism and Transformers
- Autoencoders
- Generative models: Variational Autoencoders, GAN Networks
- Graph Convolutional Networks
- *Deep Learning for Reinforcement Learning (?)*

# Course Organization

---

## Lectures:

Gorlaeus 01, Fridays, 15:15-17:00 or 17:15-19:00

<https://www.student.universiteitleid.nl/en/study--studying/study/educational-information/schedules/science/computer-science-msc?cd=computer-science-msc&cf=science#tab-2>

## Computer Labs (with own laptops(?)):

Wednesdays, 13:15-15:00

<https://www.student.universiteitleid.nl/en/study--studying/study/educational-information/schedules/science/computer-science-msc?cd=computer-science-msc&cf=science#tab-2>

*Presence during lectures & comp. lab is **not compulsory***


# Grading

---

- Course workload: 6 ects (about 168 hours)
- Final grade is based on two components:
  - Exam (40% of the final grade):  
various questions/small problems related to theory and practice (most likely via the ANS system)
  - Practicals (60% of the final grade): work in 3-persons teams  
two programming & reporting assignments (30%+30%)
  - Both parts (exam+prac) must be passed (grade > 5.5)!

# Textbook (practical)

3<sup>d</sup> edition published soon (?)




# Additional (for interested students)

---

[www.deeplearningbook.org](http://www.deeplearningbook.org)

Available on-line; print it to PDF  
or find on the internet a pdf dump;

Highly theoretical ...



# Additional Internet Resources

---

<http://deeplearning.net> (a bit old)

Blogs, Medium/TDS, Github, Neurips.cc, ...

Specific links provided during the course

**A free experimental Google platform:**

[colab.research.google.com](https://colab.research.google.com)

*Jupyter notebook server powered by CPUs, GPUs or TPUs*



# Neural Networks:

motivation

history

context

# How could we program computers to do:

---

- Optical Character Recognition
- Text2Speech translation
- Speech2Text translation
- Language2Language translation
- Robot arm control
- Face recognition
- Driving a car
- Detecting fraud with credit cards
- Playing arcade games, poker, chess, go, ...



How people do it? “training data” + “self-learning”

# The biological model

Human brain

around  $10^{11}$  neurons, and  $10^{15}$  interconnections (synapses)

## Structure of a Typical Neuron




# The biological model

## How a neuron works ?

- It receives input signals through its dendrites
- The input signals generate difference of electrical potential on the cell membrane
- The difference is propagated to the axon hillock
- A train of electrical impulses is generated along the axon
- The impulses on the axon generate the release in the synaptic space of some neurotransmitters

## Structure of a Typical Neuron



***Learning by changing  
the topology & thickness  
of connections***

# Basic principle

Artificial neural network (ANN)= set of highly interconnected simple processing units (also called neurons)

Processing unit = simplified model of a neuron

Learning = finding the architecture of connections and their weights

# Computer versus Brain

## Von Neumann machine

- One or a few high speed (ns) processors with considerable computing power
- One or a few shared high speed buses for communication
- Sequential memory access by address
- Problem-solving knowledge is separated from the computing component
- Hard to be adaptive

## Human Brain

- Large number ( $10^{11}$ ) of low speed processors (ms) with limited computing power
- Large number ( $10^{15}$ ) of low speed connections
- Content addressable recall (CAM)
- Problem-solving knowledge resides in the connectivity of neurons
- ***Adaptation by changing the connectivity (topology & thickness of connections)***

# Human Brain

---

- **graceful degradation:**  
its performance degrades gracefully under partial damage
- **self-learning:**  
it can learn (reorganize itself) from experience
- **self-healing:**  
partial recovery from damage is possible
- **massive parallelism:**  
performs parallel computations extremely efficiently. E.g., complex visual perception takes less than 100 ms, that is, 10 processing steps!
- **intelligence and self-awareness** (whatever it means)

# ANN History: Roots ...

Roots of Neural Networks are in:

**Neurobiological studies** (more than one century ago):

- How do nerves behave when stimulated by different magnitudes of electric current? Is there a minimal threshold needed for nerves to be activated? Given that no single nerve cell is long enough, how do different nerve cells communicate among each other?

**Psychological studies:**

- How do animals learn, forget, recognize and perform other types of tasks?

**Psycho-physical experiments**

- Understand how individual neurons and groups of neurons work.



# ANN History: First steps

## ❑ **Pitts & McCulloch (1943)**

- ❑ First mathematical model of biological neurons
- ❑ All Boolean operations can be implemented by these neuron-like nodes (with different threshold and excitatory/inhibitory connections).
- ❑ Competitor to Von Neumann model for general purpose computing device
- ❑ Origin of automata theory.

## ❑ **Hebb (1949)**

- ❑ Hebbian rule of learning: increase the connection strength between neurons  $i$  and  $j$  whenever both  $i$  and  $j$  are activated.
- ❑ Or increase the connection strength between nodes  $i$  and  $j$  whenever both nodes are simultaneously ON or OFF.

# ANN History: Perceptron

## ❑ Early booming (50's – early 60's)

- Rosenblatt (1958)
  - Perceptron: network of threshold nodes for pattern classification. Perceptron learning rule – first learning algorithm
  - Perceptron convergence theorem:  
*“everything that can be represented by a perceptron can be learned”*
- Widrow and Hoff (1960, 1962)
  - Learning rule that is based on minimization methods
- Minsky's attempt to build a general purpose machine with Pitts/McCulloch units

# ANN History: Setback ...

- **The setback (mid 60's – late 70's)**
  - Minsky and Papert publish a book "Perceptrons" (1969):
    - Single layer perceptron **cannot represent** (learn) simple functions such as **XOR**
    - Multi-layer of non-linear units may have greater power but there was no learning algorithm for such nets
    - Scaling problem: connection weights may grow infinitely
  - ***US Defense/Government stop funding research on ANN***


# ANN History: Backpropagation

- **Renewed enthusiasm and progress (80's – 90's)**
  - New techniques
    - **Backpropagation** learning for multi-layer feed forward nets (with non-linear, differentiable node functions)
    - Physics inspired models (Hopfield net, Boltzmann machine, etc.)
    - Unsupervised learning (LVQ nets, Kohonen nets)
  - Impressive applications (character recognition, speech recognition, text-to-speech transformation, process control, associative memory, etc.)

## **But:**

- Criticism from Statisticians, Neurologists, Biologists, Ordinary Users, ...
- Lots of ad-hoc solutions, "wild creativity"
- A lot of rubbish is produced ...

# Supervised Learning



- **Training data:** inputs & targets (= desired outputs)
- **Network:** a complicated function with many parameters to tune
- **Training:** a process of tweaking the parameters to minimize the errors of predictions (outputs should be close to targets)
- **Generalization:** we want the network to perform well on data that was not used during the training process!

$$Error = \sum_{Examples} (output_i - target_i)^2$$




# Example: Recognizing Handwritten Digits

Training set: 2000 images of digits

Each image: 16x16 pixels = a vector of 256 components

The network: 256 inputs and 10 outputs

Test set: 1000 images of digits



# Backpropagation algorithm

---

- The backpropagation algorithm searches for weight values that **minimize the total error of the network**
- It consists of the **repeated application of two passes**:
  - **Forward pass**: in this step the network is activated on one example and the **error** of each neuron of the output layer is **computed**.
  - **Backward pass**: in this step the network error is used for updating the weights. Starting at the output layer, **the error is propagated backwards through the network, layer by layer**, with help of the **generalized delta rule**. Finally, all weights are updated.
- **No guarantees of convergence**  
(when learning rate too big or too small)
- In case of convergence: **local (or global) minimum**
- In practice: **try several starting configurations and learning rates**.

# Three examples of MLP



**NetTalk:** a network that reads aloud texts

**ALVINN:** a Neural network that drives a car

**Falcon:** a real-time system for detecting fraud  
with credit card transactions




# NETtalk (Sejnowski & Rosenberg, 1987)

---

- The task is to **learn to pronounce English text** from examples (text-to-speech)
- Training data: a list of **<phrase, phonetic representation>**
- **Input:** 7 consecutive characters from written text presented in a moving window that scans text
- **Output:** phoneme code giving the pronunciation of the letter at the center of the input window
- **Network topology:** 7x29 binary inputs (26 chars + punctuation marks), 80 hidden units and 26 output units (phoneme code). Sigmoid units in hidden and output layer

# NETtalk




## NETtalk (contd.)



- Training protocol: 95% accuracy on training set after 50 epochs of training by full gradient descent.  
78% accuracy on a test set
- DEC-talk: a rule-based system crafted by experts  
(a decade of efforts by many linguists)
- Functionality/Accuracy almost the same
- Try: <http://cnl.salk.edu/Media/nettalk.mp3>

# ALVINN (1989)

D.A. Pomerleau, Autonomous Land Vehicle In a Neural Network (NIPS '89)



# FALCON: A Fraud Detection System

---

<http://www.fico.com/en/Products/DMAApps/Pages/FICO-Falcon-Fraud-Manager.aspx>

Right now, leading institutions around the world trust Fair Isaac's Falcon™ Fraud Manager to protect more than **450 million active credit and debit cards**. Why? They know they'll receive regular technological advances which will allow them to stay a step ahead of new and emerging fraud types. **Protecting 65% of the world's credit card transactions, Falcon detects fraud** with pinpoint accuracy via **proven neural network models** and other proprietary predictive technologies. Debit, credit, oil and retail card issuers in numerous marketplaces rely on Falcon to detect and stop fraudulent transactions - and combat identity theft - in real time.



# Why only “shallow networks” ?

---

- In practice, "classical" multi-layer networks work fine only for a very small number of hidden layers (typically 1 or 2) - this is (was) an empirical fact ...
- Adding layers is harmful because:
  - the increased number of weights quickly leads to data overfitting (lack of generalization)
  - huge number of (bad) local minima trap the gradient descent algorithm
  - vanishing or exploding gradients (the update rule involves products of many numbers) cause additional problems

## Overfitting:

which polynomial  
better approximates  
the green line?



The more parameters  
the higher the risk  
of overfitting !

# Why do we need many layers?



- In theory, one hidden layer is sufficient to model any function with an arbitrary accuracy; however, the number of required nodes and weights grows exponentially fast
- The deeper the network the less nodes are required to model "complicated" functions
- Consecutive layers "learn" features of the training patterns; from simplest (lower layers) to more complicated (top layers)
- Visual cortex consists of about 10 layers




# An end of hype?

---

## Dominance of new techniques (90's - 2005):

- Support Vector Machines (Vapnik)
- Kernel Methods (Vapnik, Scholkopf, ...)
- Ensemble methods (Breiman, Hasti, Tibshirani, Friedman,...)
- Random Forests
- *Neural Networks lose their prominent role in the fields of Pattern Recognition and Machine Learning*
- *since 1996 no new textbooks on Neural Networks !*



# Deep Learning Revolution

---

~ 2006 : G. Hinton's group: new ideas, inventions, applications

- Restricted Boltzmann Machine as a building block for DNN
- Contrastive Divergence algorithm
- Combine supervised with unsupervised learning
- Deep Belief Networks, Stacked Auto-Encoders

2010 : AlexNet (ImageNet Challenge): start of the DL-revolution!

2013: DeepMind “solves” the Atari 2600 games: -> Reinforcement Learning

2016-: AlphaGo, AlphaGoZero, AlphaZero (Go, Chess, ...?)

Many spectacular applications beating existing approaches

1989 (!) : LeCun et al: A Convolutional Network for OCR (AT&T) !???

2010-... : deeper networks, complicated architectures, more applications, ...

# Enabling Factors

---

- **Availability of “Big Data”:** the more data we have the better we can train a network
- **Powerful Hardware (GPU’s):** speedup of the training process by 100-1000 times reduces the training time from years to hours
- **New algorithms and architectures:** leaving the MLP standard behind ...
- ***Many spectacular successes!***

# Key Deep Learning Architectures



- **Convolutional Networks:** when adding layers enforce hidden nodes to learn "local features" - that reduces the number of parameters
- **Recurrent Networks, Transformers:** networks for modeling sequential data (text, speech, video, 3D-live brain scans, ...)
- **Autoencoders:** networks that learn intrinsic features of the data
- **VAEs, GANs:** generative models of the data
- **Graph Convolutional Networks:** networks for learning from graphs

# Autonomous Land Vehicle




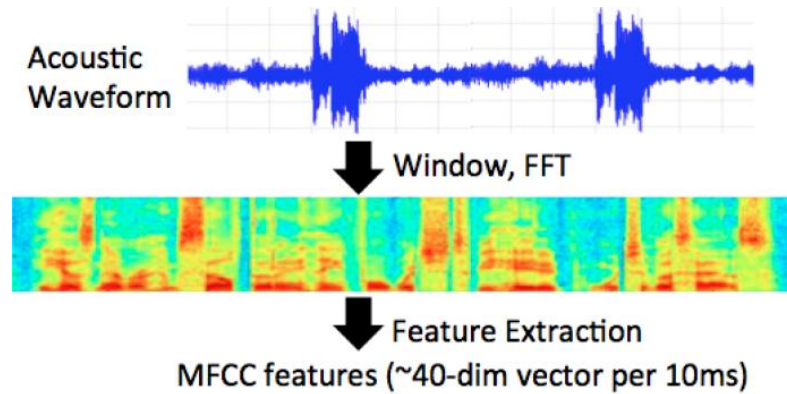
(DARPA's GrandChallenge 2005 contest)





# Other success stories

<http://cl.naist.jp/~kevinduh/a/deep2014/140121-ResearchSeminar.pdf>



# Convolution Networks

---




training:  
60.000 images

testing:  
10.000 images

each image:  
32x32 pixels

accuracy: **99.7%**  
(on the test set)

# All 33 misclassified digits





# A Convolutional Filter

We want to find locations in the that look like a 3x3 cross.

We take the matrix **F** (called **a convolutional filter**) and "convolve it" with all possible locations in the image.

We will get another (smaller) matrix with "degrees of overlap":

$$F = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

*"multiply and add"*

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature


# Filters as feature detectors

---


A filter: a “feature detector” – returns high values when the corresponding patch is similar to the filter matrix

Think about all pixels being **-1** (black) or **+1** (white) and filter parameters also restricted to **-1** and **1**

Example: what features are “detected” by:



# LeNet for digit recognition



- C1: 6 convolutional filters of size 5x5
- S2: lowers the resolution by factor 2
- C3: 16 filters of size 5x5
- ...
- The whole network has:
  - 1256 nodes
  - 64.660 connections
  - **9.760 trainable parameters (and not millions!)**
  - **trained with the Backpropagation algorithm!**

# ImageNet

- ❑ 15M images
- ❑ 22K categories
- ❑ Images collected from Web
- ❑ Human labelers (Amazon's Mechanical Turk crowd-sourcing)
- ❑ ImageNet Large Scale Visual Recognition Challenge (ILSVRC-2010)
  - 1K categories
  - 1.2M training images (~1000 per category)
  - 50,000 validation images
  - 150,000 testing images
- ❑ RGB images
- ❑ Variable-resolution, but this architecture scales them to 256x256 size

# From LeNet to ImageNet (2010/2012)

---

## ImageNet

- 15M images
- 22K categories
- Images collected from Web
- RGB Images
- Variable-resolution
- Human labelers (Amazon's Mechanical Turk crowd-sourcing)

## ImageNet Large Scale Visual Recognition Challenge (ILSVRC-2010)

- 1K categories
- 1.2M training images (~1000 per category)
- 50,000 validation images
- 150,000 testing images

# ImageNet

## Classification goals:

- ❑ Make 1 guess about the label (Top-1 error)
- ❑ make 5 guesses about the label (Top-5 error)





# ImageNet


- ILSVRC-2010 test set

Model	Top-1	Top-5
<i>Sparse coding</i> [2]	47.1%	28.2%
<i>SIFT + FVs</i> [24]	45.7%	25.7%
CNN	<b>37.5%</b>	<b>17.0%</b>

- ILSVRC-2012 test set

Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
<i>SIFT + FVs</i> [7]	—	—	26.2%
1 CNN	40.7%	18.2%	—
5 CNNs	38.1%	16.4%	<b>16.4%</b>
1 CNN*	39.0%	16.6%	—
7 CNNs*	<b>36.7%</b>	15.4%	<b>15.3%</b>

# AlexNet simplified



- not shown: max-pooling and local contrast normalization applied between convolutional layers








# AlexNet: additional “tricks”




- Computations distributed over 2GPUs
- Local Contrast Normalization
- ReLU (Rectified Linear Unit) instead of sigmoid activation functions
- L2 weight normalization: punish big weights
- Data Augmentation: increase the number of training records by applying some modifications: shifts, contrasts, ...
- Dropout: when training, in every iteration disable 50% nodes (disabling weights doesn't work!)

# Results



			
<b>mite</b>	<b>container ship</b>	<b>motor scooter</b>	<b>leopard</b>
<div> <div></div> <div>mite</div> <div>black widow</div> <div>cockroach</div> <div>tick</div> <div>starfish</div> </div>	<div> <div></div> <div>container ship</div> <div>lifeboat</div> <div>amphibian</div> <div>fireboat</div> <div>drilling platform</div> </div>	<div> <div></div> <div>motor scooter</div> <div>go-kart</div> <div>moped</div> <div>bumper car</div> <div>golfcart</div> </div>	<div> <div></div> <div>leopard</div> <div>jaguar</div> <div>cheetah</div> <div>snow leopard</div> <div>Egyptian cat</div> </div>
			
<b>grille</b>	<b>mushroom</b>	<b>cherry</b>	<b>Madagascar cat</b>
<div> <div></div> <div>convertible</div> <div>grille</div> <div>pickup</div> <div>beach wagon</div> <div>fire engine</div> </div>	<div> <div></div> <div>agaric</div> <div>mushroom</div> <div>jelly fungus</div> <div>gill fungus</div> <div>dead-man's-fingers</div> </div>	<div> <div></div> <div>dalmatian</div> <div>grape</div> <div>elderberry</div> <div>ffordshire bullterrier</div> <div>currant</div> </div>	<div> <div></div> <div>squirrel monkey</div> <div>spider monkey</div> <div>titi</div> <div>indri</div> <div>howler monkey</div> </div>

# CNNs: progress



[https://medium.com/@siddharthdas\\_32104/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5](https://medium.com/@siddharthdas_32104/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5)

# CNNs: progress




An Analysis of Deep Neural Network Models for Practical Applications, 2017.

[https://medium.com/@siddharthdas\\_32104/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5](https://medium.com/@siddharthdas_32104/cnns-architectures-lenet-alexnet-vgg-googlenet-resnet-and-more-666091488df5)

# Residual Networks

(He, Zhang, Ren, Sun, 2015)

- Key idea: it's easier to learn “the modification of the original image than the modified image”
- Implementation of the key idea: add *identity shortcuts* between 2 (or more) layers
- Huge jump in accuracy
- 1000 layers is possible!



# Reinforcement Learning & Deep Mind



**2013: DeepMind publishes a paper:**  
**DNQ: Playing Atari with Deep Reinforcement Learning**

**“zero-knowledge”:**

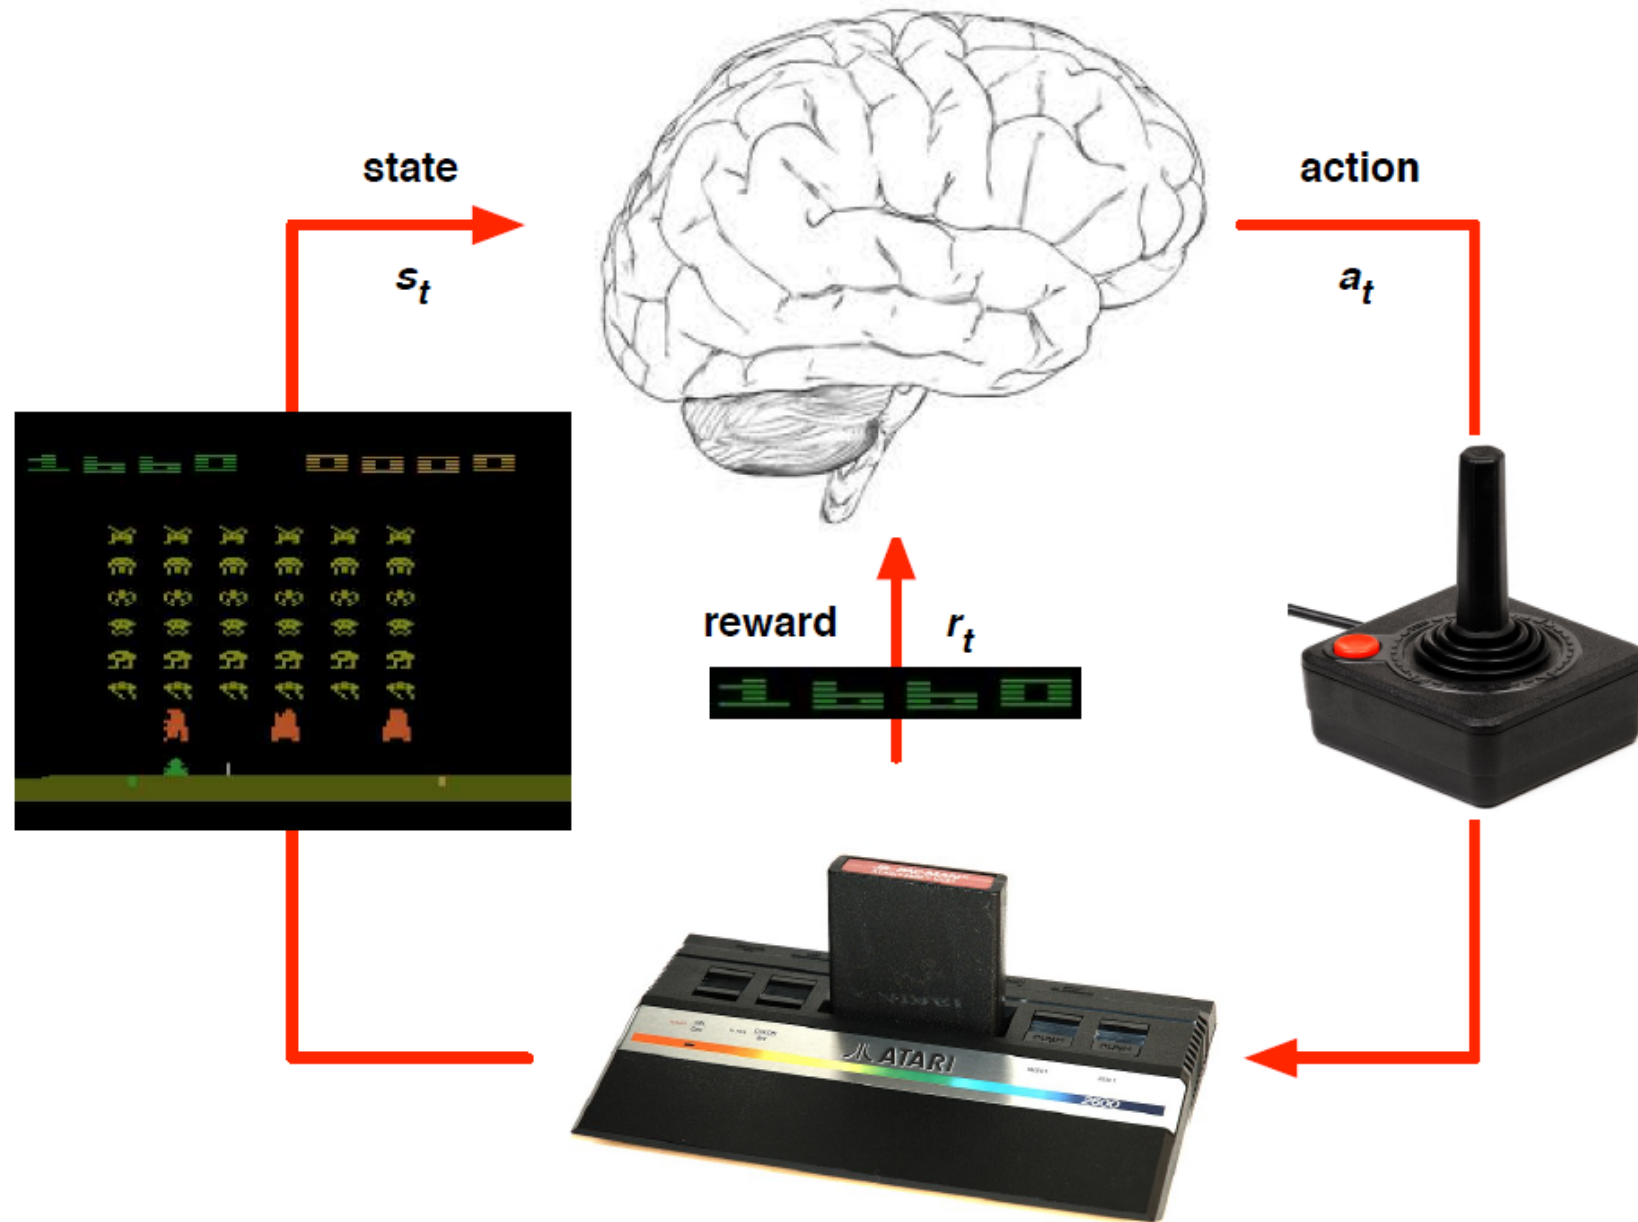
- a deep neural network can learn to **play any Atari 2600 games** using as input screen dumps and as output joystick positions
- training process: **a few months of CPU-time; 2 days of GPU-time**

**2016: DeepMind/Google develops another network**  
**AlphaGo network beats the world GO champion 6:0**

**2017: DeepMind develops a network that learns “from scratch” to play chess and GO beating anything that exists! AlphaGo-Zero**




# Reinforcement Learning in Atari



# DQN in Atari

- ▶ End-to-end learning of values  $Q(s, a)$  from pixels  $s$
- ▶ Input state  $s$  is stack of raw pixels from last 4 frames
- ▶ Output is  $Q(s, a)$  for 18 joystick/button positions
- ▶ Reward is change in score for that step



Network architecture and hyperparameters fixed across all games  
[Mnih et al.]



# DQN for Atari 2600 Games






Fig. 4. The DQN [71]. The network takes the state—a stack of greyscale frames from the video game—and processes it with convolutional and fully connected layers, with ReLU nonlinearities in between each layer. At the final layer, the network outputs a discrete action, which corresponds to one of the possible control inputs for the game. Given the current state and chosen action, the game returns a new score. The DQN uses the reward—the difference between the new score and the previous one—to learn from its decision. More precisely, the reward is used to update its estimate of  $Q$ , and the error between its previous estimate and its new estimate is backpropagated through the network.



# AlphaGo, AlphaGo Zero, Alpha Zero

---

- **October 2015:** first version of **AlphaGo** beats **European Go Champion Fan Hui 5:0**
- **March 2016:** **AlphaGo** beats **World Go Champion, Lee Sedol**  
AlphaGo was pre-trained on a huge database of historical games, and then improved by playing against itself; a few **TPUs**
- **19 October 2017:** **AlphaGo Zero**, trained solely on self-played games **surpasses AlphaGo, beating it 100:0**, hardware costs estimated on \$25 million
- **December 5, 2017:** **Alpha Zero** announced: a generic architecture learning to play **Go, Chess, Shogi in hours** (4 hours to learn Chess on a superhuman level); hardware: **5000 TPUs**

**What NEXT?**